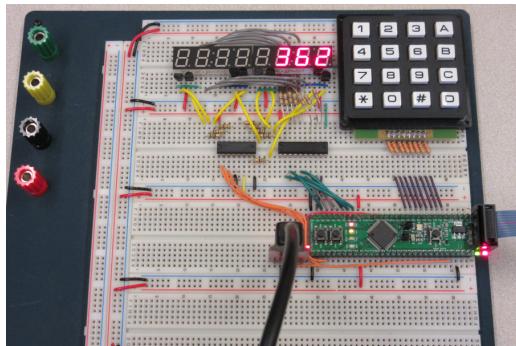


# ECE 362 Lab

## Experiment 5: Timers

[Home](#)    [About](#)    [References](#)    [Notes](#)  
[Homework](#)    [Labs](#)

### Introduction



Each timer subsystem allows the microcontroller to produce periodic interrupts while the CPU is busy with other operations. In this experiment, we will use timer

interrupts to periodically scan a matrix keypad and reliably report button presses. By recording the history of button presses with an interrupt handler, the main program is free to do other things, or it can wait for a button press to take place.

### Instructional Objectives

- Learn how to configure the timer subsystem
- Use timer interrupts
- Reliably debounce a *matrix* of buttons

- Learn how to drive a *multiplexed* display

## Table of Contents

Step	Description	Points
0	<a href="#">Prelab Exercises</a>	15
1	<a href="#">Background and wiring</a>	

\* All the points for this lab depend on proper completion of and submission of your post-lab results.

[When you are ready for your lab evaluation, review this checklist.](#)

## Step 0: Prelab Exercises:

Do them in this order...

1. Be familiar with lectures up to and including Debouncing and Multiplexing.
2. Chapter 15 of your textbook is the most helpful thing you can read to understand the timer subsystem.
3. Read though Chapter 20 of the [STM32F0 Family Reference Manual](#) to become familiar with the basic timer (TIM6/TIM7) subsystem.
4. Read the datasheets for the hardware you will use:
  - The [TDCR1050M 4-digit common-anode 7-segment display](#)
  - The [TLC59211 sink driver](#)
  - The [74HC138 3-to-8 decoder with active-low outputs](#)

5. Read this entire lab document.
6. Ensure that your development board and serial port are working **before** you start the lab.
7. After doing the previous steps, including reading the entire lab document, then do the [prelab exercises](#) and submit them **before** attempting the lab experiment.
8. Wire the keypad and 7-segment displays as described in section 1.5 of this lab document.

## Step 1: Background and wiring

### 1.1 Debouncing buttons

Because switches are mechanical devices, they bounce. This means we do not get a stable electrical signal on a press or a release. Because of switch bounce, the voltage fluctuates between  $V_{DD}$  and ground before settling down to a stable value. There are several ways to fix this.

Debouncing can be performed in hardware or software.

**1) Hardware debouncing:** You have already seen this working well in previous labs (lab 3). Hardware debouncing is done via a simple RC network to filter high frequency glitches that occur from mechanical bounce. The Schmitt trigger inputs of the microcontroller can tolerate the relatively low speed of the signal changes.

**2) Software debouncing:** In certain applications it might be beneficial to debounce switches in software rather than use an RC filter on each switch.

Software debouncing can be achieved in a variety of methods. The simplest debouncing method is to read the voltage, wait for the switch to settle and read it again. However there are more robust approaches, which are more suited for a keypad matrix. Here we describe one such method that we dub the "history method".

In this "history" method we store N previous button states. Each button's history is represented by a single variable. This variable is referred from here on, as the 'history variable'. The least significant bit of the history variable represents the most recent button state and the most significant bit represents the oldest button state in history. When a new button state is read, the history variable is shifted left by one, making space for the new state in the least significant bit. The new state value (i.e. 0 or 1) is ORed into the history variable.

Now that we have a history of button states we can look for two specific patterns to identify a button press and release. An 8-bit pattern of 00000001 represents the first moment of a button press and 11111110 represents the first moment that of a button release. Note that these patterns are applicable when the button's logic is active high, else the patterns are switched. Therefore, by checking if the history variable equals 1 (i.e. 00000001) we can detect a button press. Similarly by checking if history variable is 0xfe (i.e. 11111110) we can detect a button release.

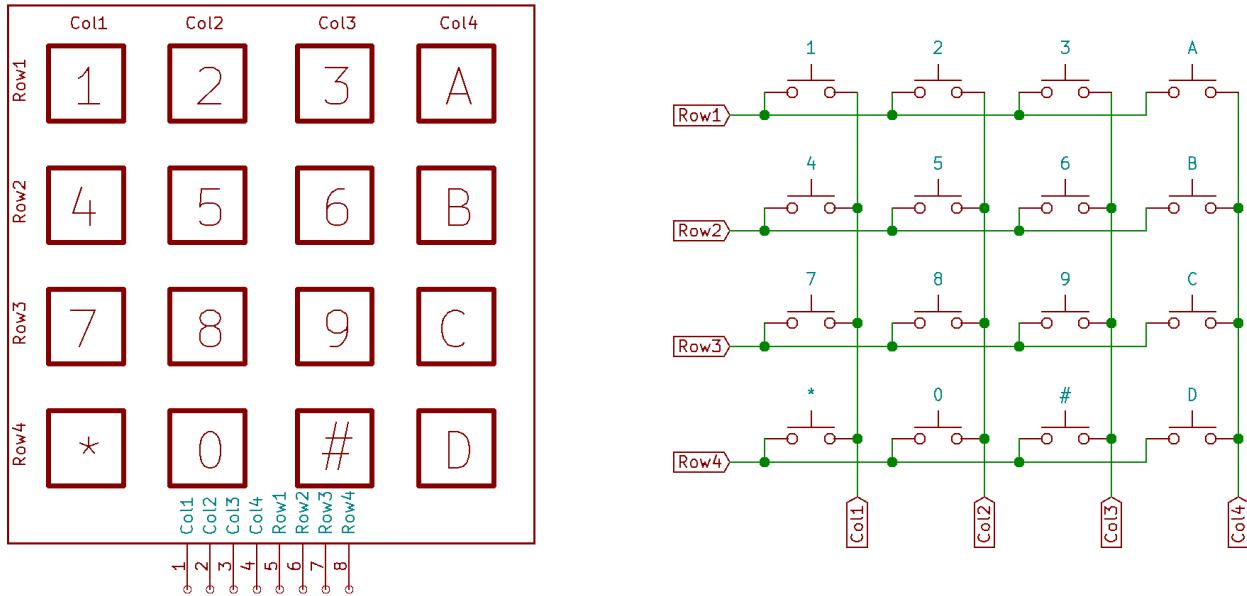
## 1.2 Reading a button matrix

A matrix of buttons presents an additional complication to the challenge of debouncing since it is not possible to monitor each button separately. To understand why this is so, it's helpful to discuss how a matrix keypad works.

You will use a 16-button keypad for this lab experiment. One configuration for such an arrangement of buttons would be to have two wires for each button (32 wires total) so that each button could be independently monitored. That would require a lot of pins on the keypad which might make it difficult to use on a breadboard. An optimization might be to have one *common* power connector and one more connector per button (17 wires total) so that each button could still be monitored individually. But this is still too many pins on the keypad than manufacturers will usually support, and the numbers would only be much worse for keypads with even more buttons.

The typical concession for keypads is to have multiple *common* pins that connect to groups of buttons. The logical extension to this trend is to configure a keypad as a *matrix* of rows and columns. If the number of rows and columns is about equal, this practice minimizes the number of pins on the keypad. The added challenge is that the keypad must be scanned one row or column at a time. For instance, a voltage can be applied to a column of buttons, and the rows can be sensed to detect if any of the buttons in the column are being pressed. Then the voltage can be removed from the column and re-applied to another column. By cycling through all the columns, all keys can eventually be detected. The difficulty is that, not only do the

buttons still bounce, the bouncing must be tracked over multiple successive scans.



**Figure 1: Keypad schematic (click for larger view)**

The schematic for the keypad matrix you will use for this lab experiment is shown in Figure 1. For this experiment, you will implement the 'history' method of debouncing buttons. For each of the 16 keys, an 8-bit variable is used. Each history variable is stored as an element of a 16-byte array in RAM. The array's first 4 elements, 0 – 3 represent the history for the first column of keys, the elements from 4 – 7 represent the history values of the second column of keys and so on. Since buttons will be sampled four at a time, the history values will also be updated four at a time. This is effectively equivalent to sampling a single button at a time except that the samples are further apart. Button presses and releases are still detected the same way.

Another minor downside of this method is that it is not always possible to detect multiple simultaneous button presses. For instance, if Col1 has an applied voltage, it is possible to detect the '1' and '4' buttons pressed simultaneously because they will be connected to distinct rows. It is also possible to detect two buttons in different rows and different columns. Nevertheless, if two buttons in the same row are pressed simultaneously, the results will be unreliable. For instance, if '1' and '2' are pressed, when Col1 has a high voltage applied, Col2 will have a low voltage applied, and vice-versa. The value read on the Row1 will be indefinite. There is a way to overcome this difficulty: use pull-up resistors on the columns as well as the rows, and configure the driven pins to use open-drain outputs so that they can only be pulled low. Then, the presence of a high voltage read on the row pins always indicates "no button in this row pressed" and a low voltage on one or more row pins indicates a button press. (Inverted thinking, like this, is difficult when programming in assembly language. We might give this a try in the future, but not now.)



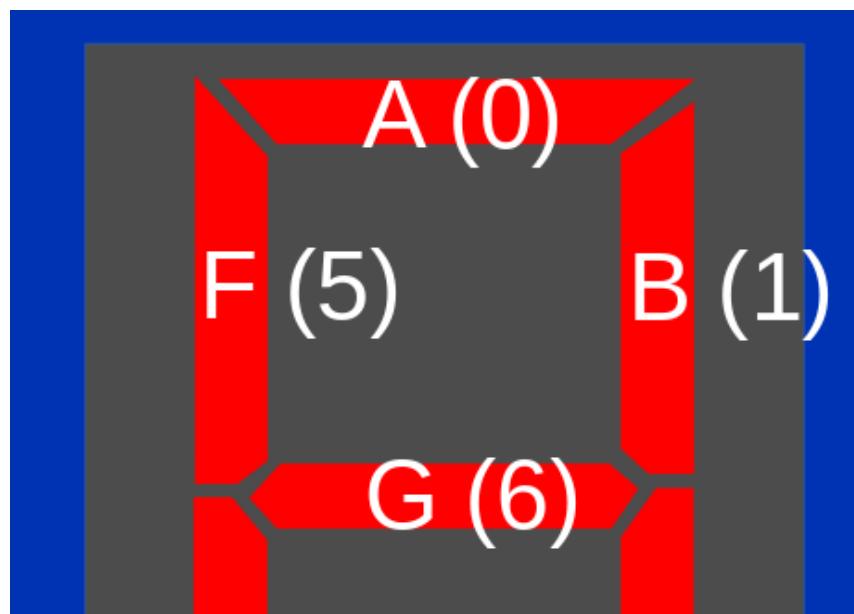
In the event that logic high and logic low are applied in a way that current flows through a button, the keypad button may be damaged. In this experiment, you will be driving every column either high or low. If two buttons are pressed in the same row, that could mean lots of current flowing from one column to another. We don't know the current

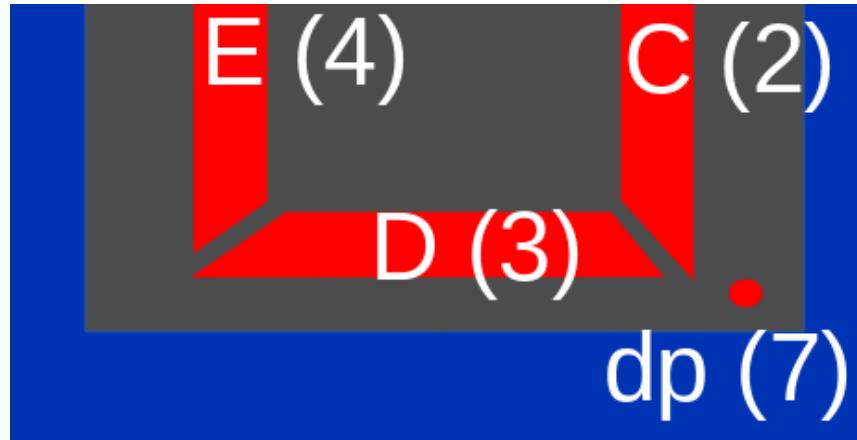
rating of the keypad switches, but we know they

weren't made for this. For this reason, you will use  $1\text{K}\Omega$  series resistors between the STM32 and all of the pins of the keypad. By doing so, there cannot possibly be any damage.

### 1.3 Driving multiplexed displays

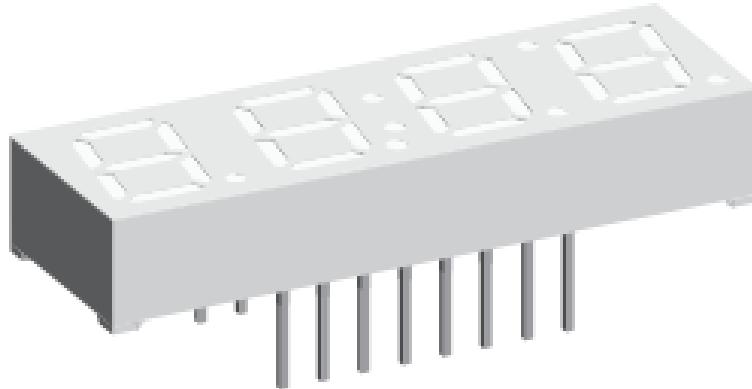
A 7-segment display is one that has seven LED segments in the shape of a number "8". By selectively turning segments on, other numbers can be displayed. In practice, such displays have more than seven segments. An eighth segment shows a decimal point. Recall Figure 2, from ECE 270, that shows the names of the segments of a display. The top segment is customarily called "A". Proceeding clockwise, the next ones are named "B", "C", and so on. The middle segment is named "G". The decimal point is named "DP".





**Figure 2: Seven-segment display organization**

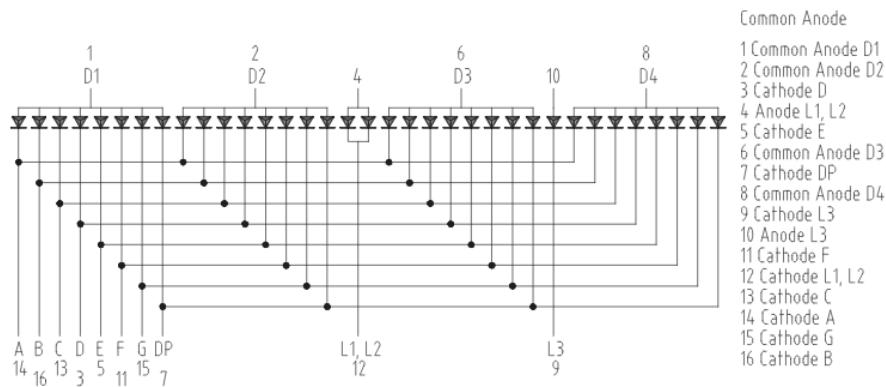
For this lab experiment, you will be using the TDCR1050M 4-digit 7-segment LED display. This device was made to be used in a clock, so it has a middle colon for use in displaying a time. Another dot on the display can be illuminated to indicate things like an alarm setting. For this experiment, we will use only the digits and decimal points.



**Figure 3: Picture of physical display**

If you look at a picture of the physical device in Figure 3, you see that it has only two rows of eight pins on each side. Only sixteen pins are used to control more than 30 individual LED segments that must be independently lit to form four distinct digits and decimal points. To do so, it

uses a "common anode" configuration so that one pin is connected to the anodes of all the LEDs on one digit. There are four distinct pins connected to the anodes of the four digits. There are another eight pins connected to similar cathodes on all the displays. For instance, one pin is connected to the four "A" cathodes on the four digits. This means it is possible to illuminate all four of the "A" segments on the four digits by applying a positive voltage to the four common anodes and a negative voltage to the pin connected to all of the "A" segments. It is not possible to simultaneously illuminate the "A" segment of all the displays and illuminate the "B" segment of only a single display. Figure 4 shows the schematic for the connections to all of the LED segments for the display.



**Figure 4: Schematic for the TDCR1050M display**

If it is not possible to independently turn on and off individual segments at the same time, how is it possible to display distinct numerals on the four groups of segments? This display is made to be *multiplexed*. In a manner similar to how we cannot detect the exact moment that a button in a matrix is pressed, we cannot simultaneously turn on

every LED segment we need. Instead, the segments of the first digit are turned on, then the segments of the next digit are turned on, then the third, then the fourth. If each individual digit is illuminated for one second, then the display will not be perceived as a four-digit number. If, however, the cycling through the digits is done quickly enough, it will *appear* that all of the numbers are shown simultaneously. This is because of the human eye's *persistence of vision*. A flickering light—as long as the flickering is rapid enough—is perceived as continually on, although maybe dimmer than normal.

As an example, consider the result of applying the following cycle of voltages applied to the pins of the display:

- A positive voltage to pin 1 (D1), and a negative voltage to pins 16 (B) and 13 (C)
- A positive voltage to pin 2 (D2), and a negative voltage to pins 14 (A), 16 (B), 3 (D), 5 (E), and 15 (G)
- A positive voltage to pin 6 (D3), and a negative voltage to pins 14 (A), 16 (B), 13 (C), 3 (D), and 15 (G)
- A positive voltage to pin 8 (D4), and a negative voltage to pins 16 (B), 13 (C), 11 (F), and 15 (G).

When these four configurations are applied repeatedly and at high speed, it will *appear* that the four-digit number "1234" is present on the display.

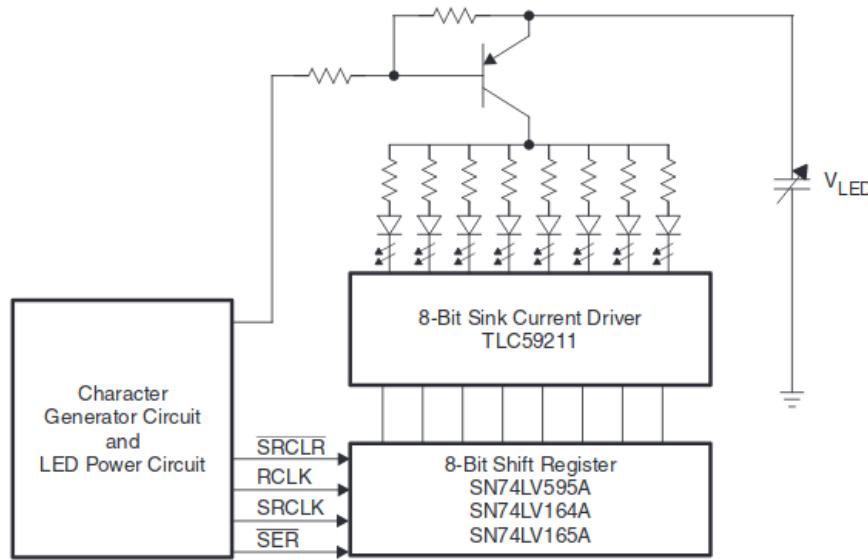
How rapidly must you cycle through each digit to achieve the effect of persistence of vision? Usually, anything flickering faster than 30 times per second appears to be continually on. As long as all of the digits of the display are

turned on and off more than 30 times per second (even though none of them are on at the same time as others) the display will appear continuous. In this lab experiment, you will use two groups of the four-digit displays (eight digits in total), and you will turn on each digit for one thousandth of a second. By illuminating each individual digit for one-thousandth of a second, any one of the eight displays will be illuminated eight times less frequently (or 1/125th of a second). This is much faster than is needed to perceive the digits as all being on simultaneously.

## 1.4 Use of a sink driver and high-side transistors

Most students, when they approach the idea of multiplexing, would be more comfortable with a "common cathode" display instead since it could be configured so that a positive voltage applied to the pin connected to the anodes of all "A" segments would illuminate that segment when a negative voltage is applied to the pin connected to all the cathodes of the digit. If you were connecting the pins directly to your microcontroller, that might even make sense. Nevertheless, the current-flow requirements for the LEDs make it impractical to illuminate many of them using only the capabilities of the microcontroller. Instead, it is common use *sink driver* connected to the cathodes of each type of segment (e.g., "A", "B", "C", etc...). A sink driver is a device that is capable of sinking (connecting to ground) a connection with a large current flow.  $47\Omega$  limiting resistors will be placed in between the sink driver and each cathode to prevent too much current from flowing through each

segment. Conveniently, the sink driver will sink current through an output pin when its corresponding input pin is high. In this way, it acts as an open-drain inverter. (A sink driver cannot push an output pin high.) Using the sink driver, a logic high applied to the driver will cause the particular segment to illuminate, which is how you would like to think about it.



**Figure 5: Typical LED display circuit**

Consider the circuit in Figure 5 that appears in the TLC59211 datasheet. In this lab experiment, you will wire the microcontroller to the sink driver rather than an 8-bit shift register. The device at the top of the diagram is a PNP bipolar junction transistor (BJT). This device will conduct current through the top pin (the emitter) in the direction of the arrow, and out the bottom pin (the collector) when the middle left pin (the base) has a voltage lower than the emitter (and a small current coming out of it). It is conceptually similar to a P-Channel MOSFET except that a MOSFET does not require a substantial current flow

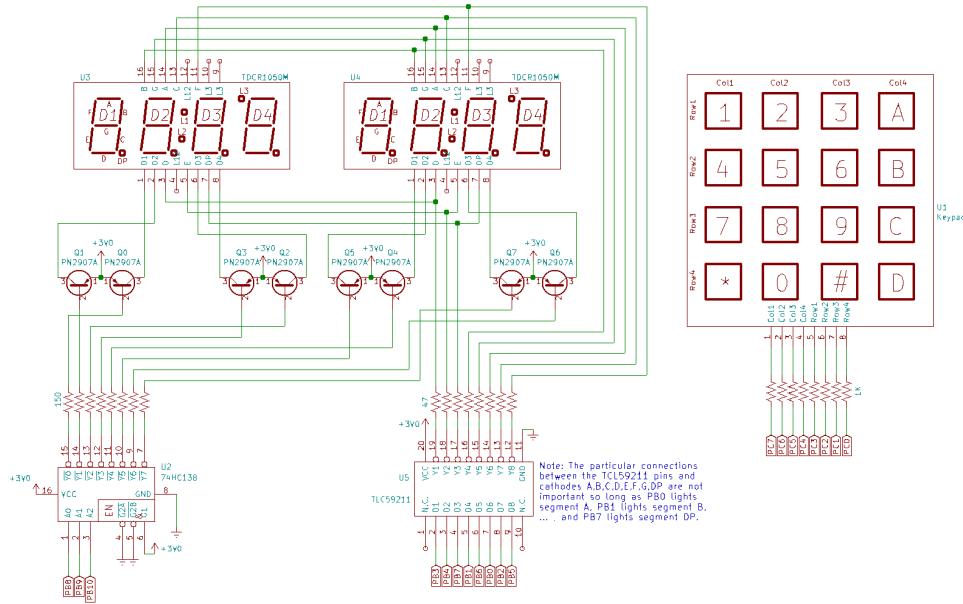
through its gate. By contrast, the collector current for a BJT is directly proportional to its base current.



You will use one PN2907A PNP BJT (in a TO-92 package) for each anode of each digit of the display. Pay careful attention to the designations for the pins (E,B,C for emitter, base, collector) on the device. Pin 1 is the emitter, pin 2—in the middle—is the base, and pin 3 is the collector. The anode of a digit must be able to supply enough current for all eight segments of the digit simultaneously. The datasheet for the TDCR1050M shows that each of the eight segments of each digit could use as much as 25 mA, although you won't have that much current flow for this experiment. A microcontroller pin would not be able to supply enough current to illuminate all of the segments of a digit, but a PN2907 transistor will provide at least 200mA if needed. Because the transistor is driving the positive voltage to the display, it is referred to as the "high-side" driver for the display. You may think of the sink driver as the "low-side" driver for the display.

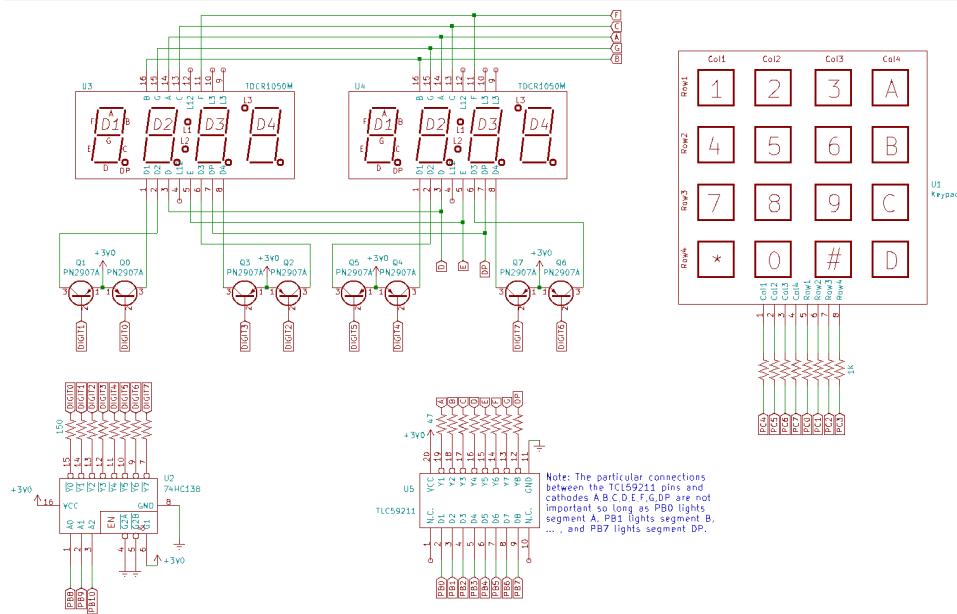
Rather than try to deal with low voltage to enable current flow through each PNP transistor, we will add one more device: a 3-to-8 decoder with active-low outputs. Three input pins select one of eight outputs to be driven low. Each one of these outputs will be connected, through a series resistor, to the base of one high-side driver transistor. Altogether, eleven pins of the microcontroller will be used to drive the display: three to select the digit to turn on through

the decoder and high-side driver transistors, and eight more to turn on the segments of the selected digit. A full schematic of the wiring for the display and keypad is shown in Figure 6.



**Figure 6: Schematic for lab experiment wiring**

It can be difficult to follow all of the connections as they cross over each other. A second schematic, with most of the wires replaced by symbolic connections is shown in Figure 7. That should offer slightly more clarity in how the connections are made. A schematic is normally intended for a high-level understanding of the logical connections, and has few details about the physical layout of components. Nevertheless, some nuances of the physical construction are preserved in these schematics. In particular, note that the transistors are numbered, from left to right, Q1, Q0, Q3, Q2, and so on. This is intentional since that is the way they will be placed on the breadboard.

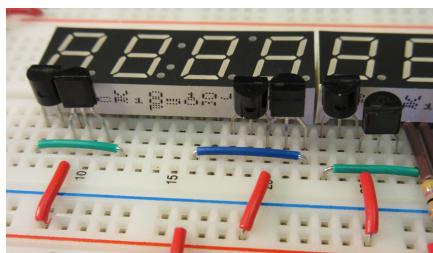


**Figure 7: Simplified schematic for lab experiment wiring**

## 1.5 Tips for wiring the circuit

You'll be assembling two 4-digit display modules, driven by 8 PN2970A transistors, a TLC59211 sink driver, 8 resistors between the sink driver and LED cathodes, a 3-to-8 decoder, and 8 more resistors between the decoder and the transistors. Another eight resistors will be used to connect the keypad to the microcontroller. With some planning, it will be possible to assemble this quickly and reliably. Here are several suggestions to make it as easy as possible...

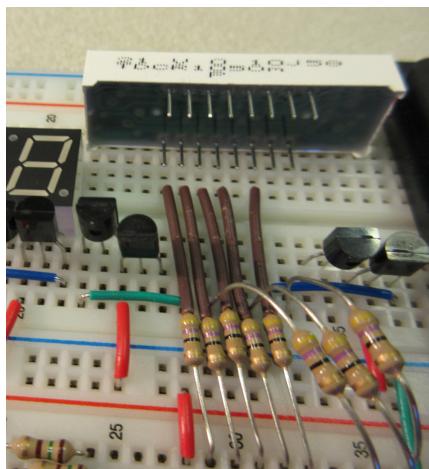
### 1.5.1 Place the transistors next to the displays



The collector of a PN2907A transistor must be connected to each of pins 1, 2, 6, and 8 of both display modules. Notice that these

pins are all on one side of the display, and that two of them are near the ends of the row of pins. There are no connections to the left or right of the connector, but there is a lot of space taken up by the displays. Rather than run long wires from the transistors to the displays, put the collector of one transistor on pin 1 and the collector of another transistor on pin 8. The other pins of each transistor are placed to the left and right, respectively. Add another two transistors further out by combining the emitters on the same breadboard row. Finally, connect the collectors of the outermost transistors to pins 2 and 6. Looking from left to right, the first backward-facing transistor will connect power to D2 (pin 2). The next forward-facing transistor will connect power to D1 (pin 1). The next transistor will connect power to D4 (pin 8), and the fourth transistor will connect power to D3 (pin 6). Note the alternating backward-facing, forward-facing manner in which the transistors are placed. The base (middle pin) of each of the transistors is what will be driven, through a  $150\Omega$  resistor, by the decoder.

### 1.5.2 Wire the resistors directly to the display

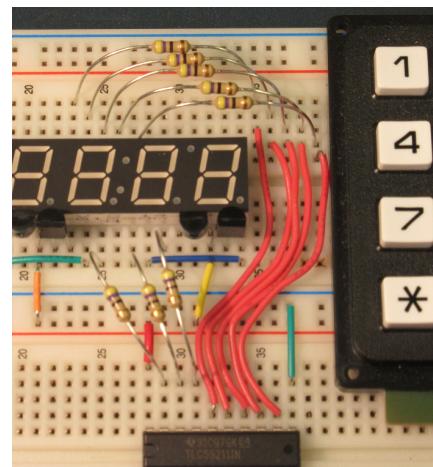
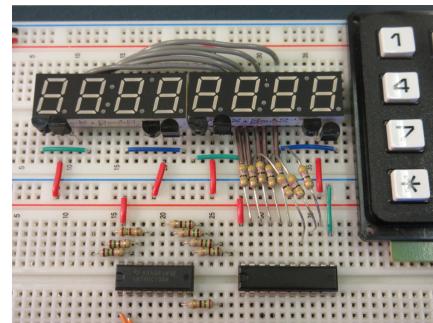


Five of the LED cathodes are on the "top" side row of connectors (pins 16,15,14,13,11) of the display, and three of them are on the "bottom" row of connectors (pins 3,5,7). Each of the eight cathodes must be connected to an output of the TLC59211 sink driver

via a series  $47\ \Omega$  resistor. If you have bulk 22 gauge copper wire, you can remove a suitable amount of insulation and put it on each of the five resistors to keep them from touching the pins they run near. Run the resistor leads between the pins below the display. Place the TLC59211 on an adjacent breadboard, in the right position to directly connect the resistors. The other three resistors are easy to wire between pins 3,5,7 and the sink driver.

Place the resistors between the cathodes and the TLC59211 outputs in the most convenient way possible. Connect the wires between the STM32 Port C and the TLC59211 so that PB0 lights segment A, PB1 lights segment B, ..., PB6 lights segment G, and PB7 lights segment DP.

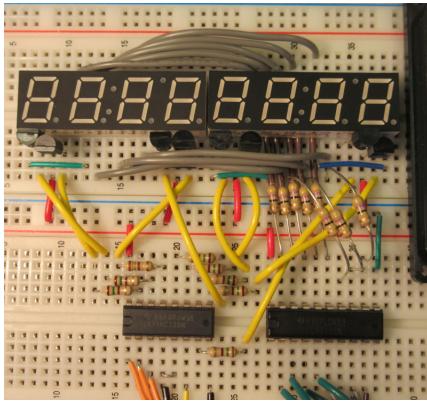
If you are not able to insulate resistor leads in this way, leave some space on the right side of the rightmost display module, connect the resistors on the top side of the display module, and run wires from the other side of the resistors to the TLC59211. It's your choice. Either way will work well. Neatness of connections will avoid a great deal of problems when you reach the point of debugging. A maze of wires that loop far over the breadboard will not only block the LED segments, it will be difficult to diagnose, and be prone to damage when the wires are snagged on objects while in transport.



### 1.5.3 Chain the cathodes from one display to the other

The eight connections to the cathodes of display on the right should be made to the same eight pins on the display on the left. Use eight wires to connect them together. I.e., connect pin 16 on the left display to pin 16 of the right display, pin 15 to 15, etc. The pins to connect together on the top side of the displays are 16, 15, 14, 13, and 11. The pins on the bottom side of the displays are 3, 5, and 7.

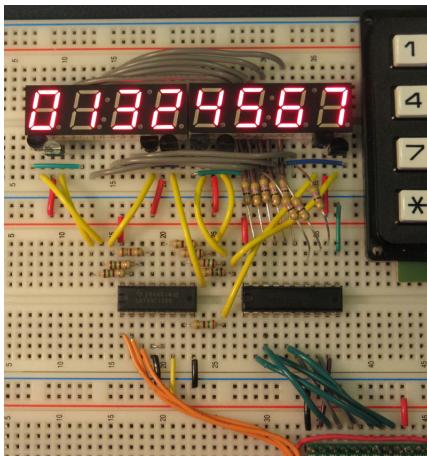
### 1.5.4 Connect resistors to the decoder



The  $150\Omega$  resistors should be placed directly on the 74HC138 decoder. If you want to be tidy, you might cut the leads to minimize the amount that they stick out of the breadboard. If you do so, make sure that leads of every resistor is wide enough to span four breadboard holes. A width of only three would mean, for instance, that the resistor on the Y2 output would connect to the  $V_{CC}$  pin of the chip. Wires should be run from the other ends of the resistors to the bases of the transistors. Connect output Y0 through a  $150\Omega$  resistor to the base of the PNP transistor connected to pin 1 of the left display. Note that output Y7 is on the bottom side of the 74HC138 decoder on pin 7. Pin 8 is the ground connection. The decoder's select lines, A0, A1, and A2 (pins 1,2, and 3, respectively) should be connected to PC6, PC7, and PC8. Pins 4, 5, and 6 should

be connected to GND, GND, and V<sub>CC</sub>, respectively. These three pins are the *enable* pins for the decoder. Unless all three of these pins are enabled, **none of the decoder outputs will be asserted** (driven low).

### 1.5.5 Check your wiring



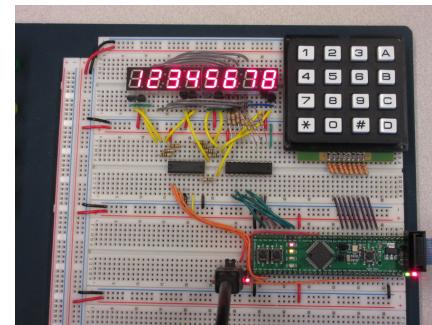
Look ahead to Section 2 of this lab document to find out how to set up the main.s template and obtain the autotest.o module. The autotest.o module for this experiment also has a wiring checker. Uncomment the `b1 check_wiring` call in main.s. When first started, the

digits 01234567 should appear on the seven-segment displays. The decimal point on the leftmost digit should illuminate for 0.25 seconds, then turn off and the decimal point on the next digit should illuminate for 0.25 seconds. The decimal point should cycle through all of the digits in two seconds and continue again with the leftmost. The picture to the left shows a situation where the bases of two transistors are connected to the wrong decoder output resistors. Instead of 01234567, the display shows 01324567. An easy mistake to make as well as to correct.

As the decimal point illuminates, make sure that the other segments of each digit remain at the same brightness. If the other segments of a digit grow dimmer, it may be that the driver transistor for that digit is *backward*. A BJT transistor normally amplifies a small base-emitter current

with a large collector-emitter current, but a BJT can also work backward so that it amplifies a small base-collector current to a larger emitter-collector current. The amplification is not as high, and it limits the current provided to a digit. The more segments per digit are lit, the dimmer they are. Carefully note the orientation of the transistors connected to the D1, D2, D3, and pins of each display.

Once the display is showing the correct digits and the decimal point is repeatedly moving from left to right, try pressing buttons on the keypad. Each button will cause a corresponding letter to be shifted onto the display to the rightmost digit, and the other seven digits will shift left. The asterisk is represented with a Greek Xi ( $\Xi$ ), and the **octothorpe** is represented with an "H". (You think I'm inventing a new word here? Go look that up.) If you press the '1' button, and a new 'A' gets shifted in on the right, it means you've switched the Row1 and Row4 wires on the keypad. If you press the '3' button, and a new '0' is shifted in on the right, it means that the wires to Row2 and Row3, as well as the wires to Col1 and Col4, are switched. By pressing buttons and looking at the outcome, you can determine which wires are incorrect.



## 2.0 Experiment

The rest of this will be published as soon as we can make it small enough that you can actually get it done this week. In

the meantime, you should get to work on the wiring.

Questions or comments about the course and/or the content of these webpages should be sent to the [Course Webmaster](#). All the materials on this site are intended solely for the use of students enrolled in ECE 362 at the Purdue University West Lafayette Campus. Downloading, copying, or reproducing any of the copyrighted materials posted on this site (documents or videos) for anything other than educational purposes is forbidden.